



**Neural networks through ice cream sales**

# Objectives:

- To *fluently* use the *vocabulary* of neural networks
- To connect *familiar past algorithms* to *neural networks*, and therefore demystify the math
- To practice *drawing* and *specifying* the parameters of a neural net
- To list the *parameters* one can *adjust* when building a neural net



# Scenario

You own a chain of ice cream stores.

You want to build a model that will predict the sales numbers of a store, given the store's location, pricing of product, and perceived quality of the product.

Simpler models haven't produced great results, so you want to try a neural network. Plus, neural networks sound fancy. You like fancy.








# Problem summary




# Variables


**Location**



**Pricing**




**Perceived Quality**




## Variables


**Location**



**Pricing**




**Perceived Quality**




## Target

**Sales**




## Variables


Location



Pricing




Perceived Quality




## Target

Sales




## Variables


Location



Pricing




Perceived Quality



**Predict**

## Target


Sales





## Variables


Location



Pricing



Perceived Quality




**Predict**

**Using a Neural  
Network**

## Target

Sales






# Vocabulary




# Using a Neural Network

## Variables


Location



Pricing




Perceived Quality



## Target


Sales




# Using a Neural Network

## Variables


Location



Pricing




Perceived Quality



input layer

## Target


Sales




# Using a Neural Network

## Variables


Location



Pricing



Perceived Quality




**input layer**

nodes: 3

## Target


Sales




# Using a Neural Network

## Variables


Location



Pricing



Perceived Quality




**input layer**

nodes: 3

## Target

Sales




**output layer**


# Using a Neural Network

## Variables


Location



Pricing



Perceived Quality




**input layer**

nodes: 3

## Target

Sales




**output layer**  
nodes: 1



## Variables


Location







Pricing



Perceived Quality




## Using a Neural Network



## Target

Sales



**input layer**


nodes: 3

**output layer**


nodes: 1

## Variables


Location



Pricing







Perceived Quality



**input layer**

nodes: 3


## Using a Neural Network



**hidden layer**

## Target

Sales




**output layer**

nodes: 1

## Variables


Location



Pricing







Perceived Quality



**input layer**

nodes: 3

## Using a Neural Network




**hidden layer**

nodes: 4

## Target

Sales



**output layer**

node: 1



**How many layers in our neural network?**



**Keep this in mind as we build neural networks**




# Draw this out:

We want to build a neural network using gender(assume binary), years of education, marital status(single vs wed), and years of employment to predict income.

We are going to use two hidden layers. The first one will have three nodes and the second will have 5.


Draw and compare w neighbors.

# The math behind networks is not that scary







# We need some notation to make this work



May have seen diagrams like this



For simplicity we are only going to focus on one layer




**input layer**

nodes: 3

**hidden layer**

nodes: 4

And specifically the first node in the hidden layer




input layer

nodes: 3


hidden layer

nodes: 4

**Problem:** this common diagram isn't representative




# What is shown as one is really three




input layer  
nodes: 3


hidden layer

Each node that receives input from previous nodes actually has **three** parts







Each node that receives input from previous nodes actually has **three** parts










Each node that receives input from previous nodes actually has **three** parts



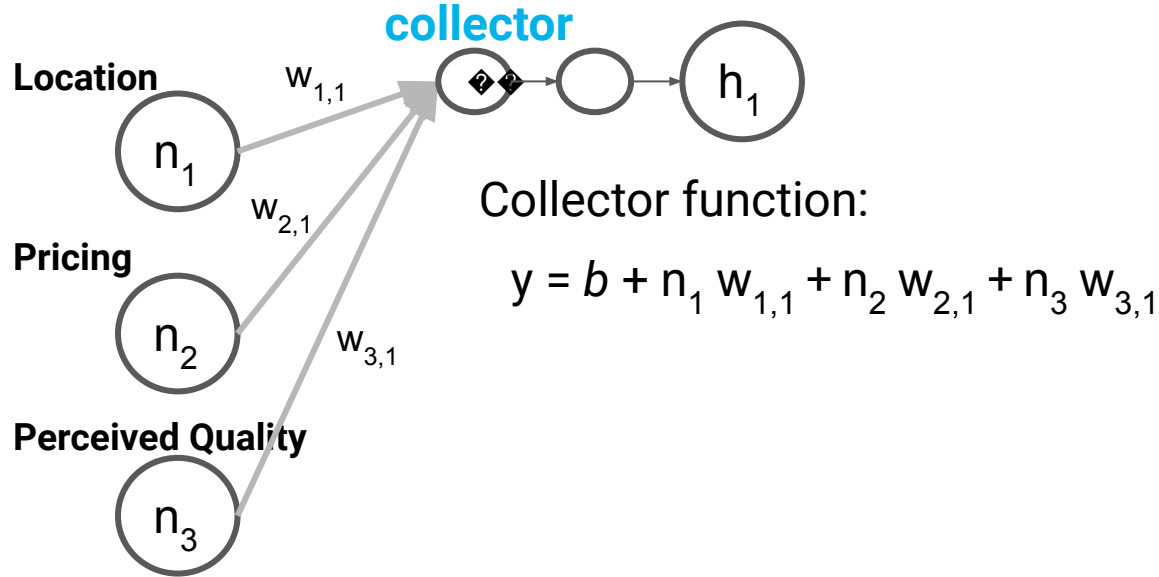


Each node that receives input from previous nodes actually has **three** parts





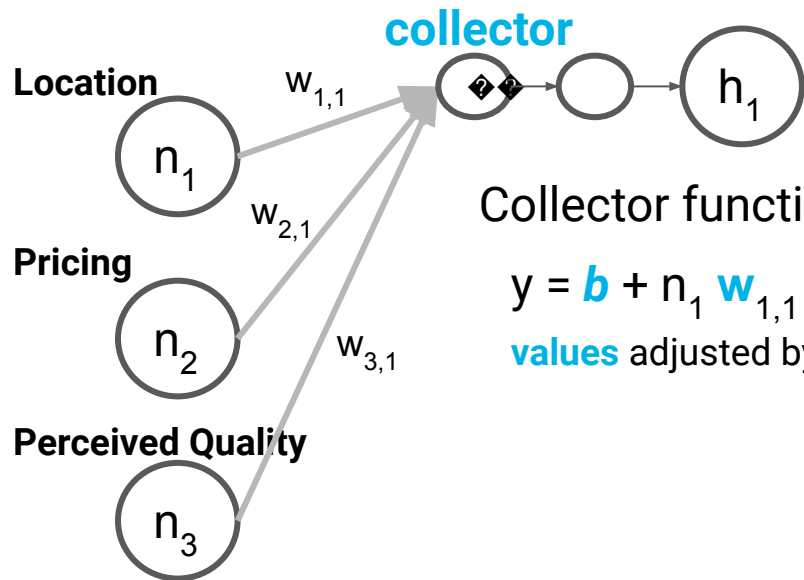
This function should look familiar



**input layer**

nodes: 3

**hidden layer**



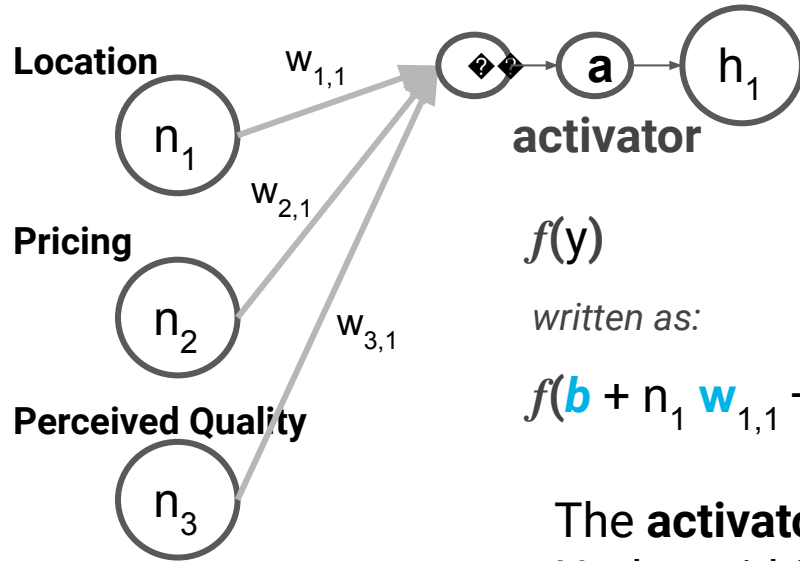
Collector function:

$$y = b + n_1 w_{1,1} + n_2 w_{2,1} + n_3 w_{3,1}$$

values adjusted by the algorithm

input layer  
nodes: 3

hidden layer



**input layer**

nodes: 3

**hidden layer**

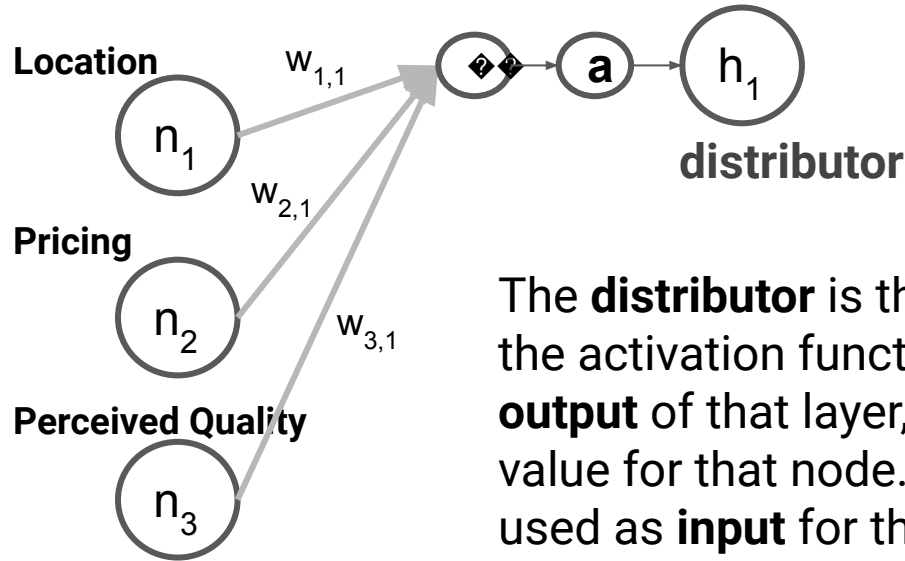
The **activator** is a function chosen by **you** that takes the output of the collector as input.

$f(y)$

*written as:*

$$f(b + n_1 \mathbf{w}_{1,1} + n_2 \mathbf{w}_{2,1} + n_3 \mathbf{w}_{3,1})$$

The **activator** is specified for **each** layer. Nodes within a layer all use the same activation function



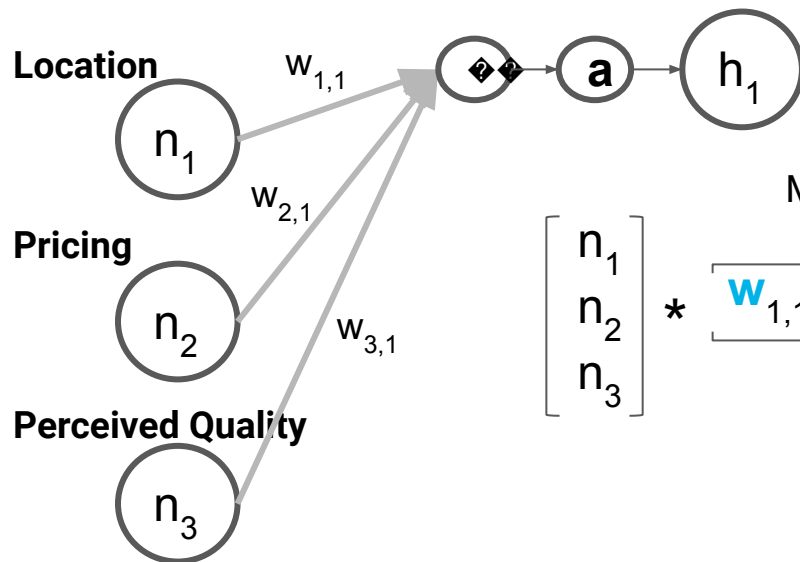
The **distributor** is the **output** of the activation function. It is the **output** of that layer, the final value for that node. It is then used as **input** for the next layer.

input layer  
nodes: 3

hidden layer



For the math lovers in the room:



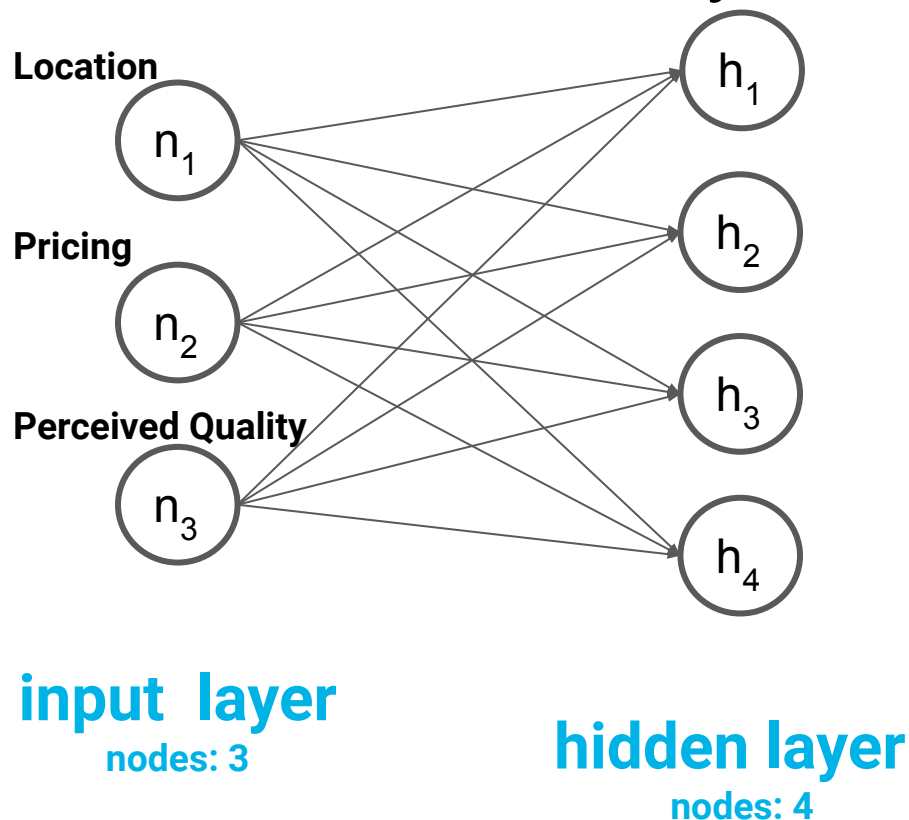
Matrix notation for one hidden node

$$\begin{bmatrix} n_1 \\ n_2 \\ n_3 \end{bmatrix} * \begin{bmatrix} \mathbf{w}_{1,1} & \mathbf{w}_{2,1} & \mathbf{w}_{3,1} \end{bmatrix} + \mathbf{b}_1 = y_1 \rightarrow f_{a1}(y_1) \rightarrow h_1$$

input layer  
nodes: 3

hidden layer

# The full math for a layer:



Matrix notation for one hidden layer

$$\begin{bmatrix} n_1 \\ n_2 \\ n_3 \end{bmatrix} * \begin{bmatrix} w_{1,1} & w_{2,1} & w_{3,1} \\ w_{1,2} & w_{2,2} & w_{3,2} \\ w_{1,3} & w_{2,3} & w_{3,3} \\ w_{1,4} & w_{2,4} & w_{3,4} \end{bmatrix}^T + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix}$$

$$\begin{bmatrix} f(y_1) \\ f(y_2) \\ f(y_3) \\ f(y_4) \end{bmatrix} = \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{bmatrix}$$




When creating a layer, what are the specs  
that **you** choose?



# Summary

Hidden layer variables	You define	Computer figures out
weights		✓
activation function	✓	
bias		✓
number of nodes	✓	



# Summary so far:


What you choose when building a neural network:

At the **network** level:

- Number of input variables
- Number of hidden layers
- If it is a classification or regression problem

At the **layer** level:

- The number of nodes
- The activation function





**What else can we adjust?**





# What you choose when building a neural network:

At the **network** level:

- Number of input variables
- Number of hidden layers
- If it is a classification or regression problem
- **Batch size**
- **Number of epochs**
- **Learning rate & optimizer**
- **Regularization type and lambda**

At the **layer** level:

- The number of nodes
- The activation function





# Batches and Epochs are about data processing

That's a lot of math and a lot of data. The dataset is split into chunks and passed through the network one chunk at a time.

**Batch** defines the number of observations in each “chunk”.


**Epochs** are how many times you want the whole dataset to go through the network.

All the **batches** = one **epoch**.




# Learning rate is from gradient descent

Scenario	Cost/Loss Function
Regression	MSE
Binary classification	Cross-Entropy (Logarithmic loss)
Multi-class classification	Softmax of Cross-Entropy




FP




# Optimizer is how the gradient is calculated

Options
sgd
rmsprop
Adagrad
Adadelat
Adam
And more!



FP



# Regularization - adjusts the weights by layer

Adding a set penalization term at each collector node.

You can choose no regularization, lasso, or ridge.

# Write in sentences what this code does

```
model = Sequential()
model.add(Dense(32, activation='relu', input_dim=100))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])


# Generate dummy data
import numpy as np
data = np.random.random((1000, 100))
labels = np.random.randint(2, size=(1000, 1))

# Train the model, iterating on the data in batches of 32 samples
model.fit(data, labels, epochs=10, batch_size=32)
```

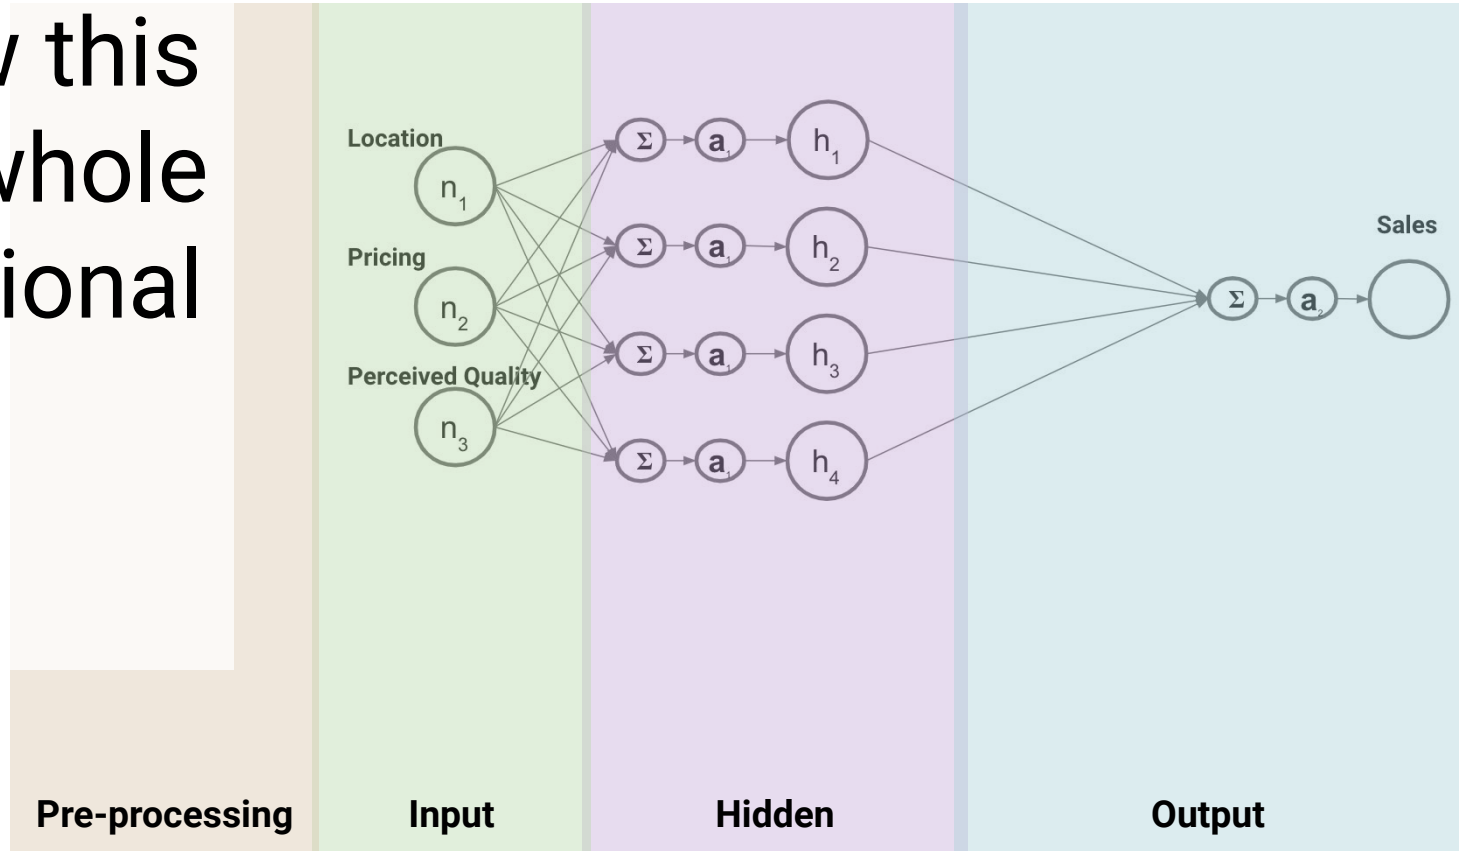
Even without learning Keras explicitly, you should be able to recognize keywords and concepts based on this review.

Dense = fully connected to all previous nodes

While we  
have  
covered  
this:



Know this  
is a whole  
additional  
area.



Exit ticket